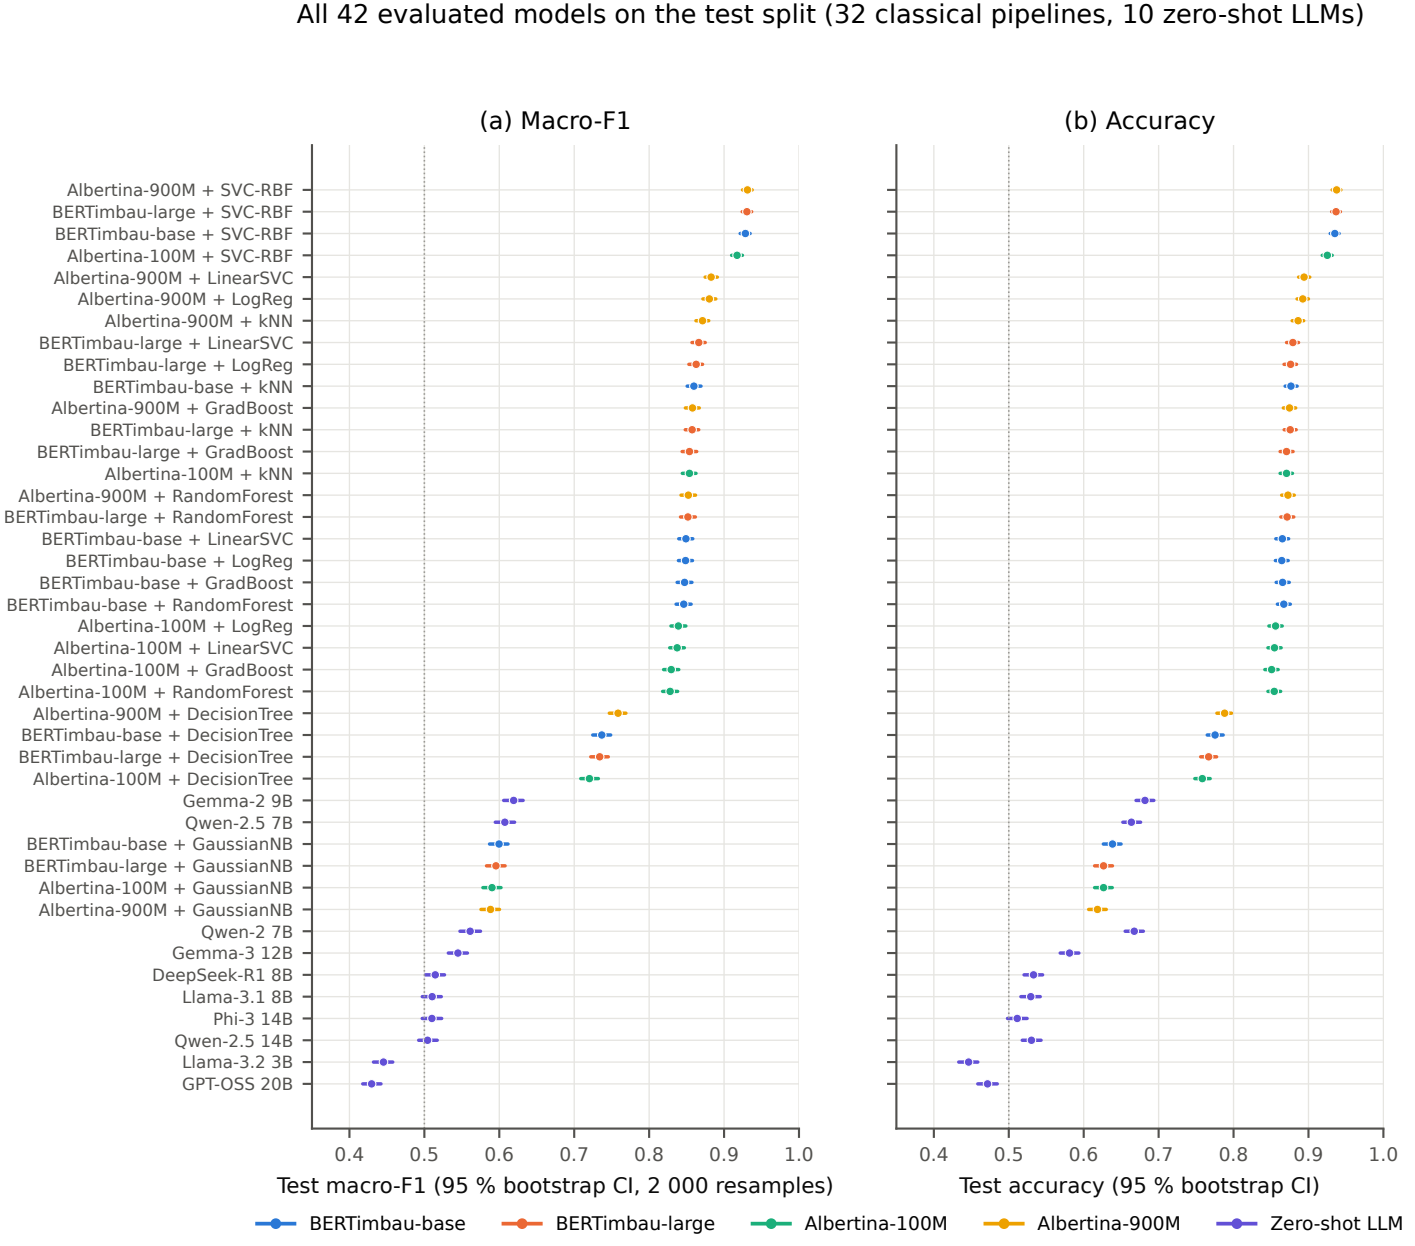
All 42 evaluated models on the test split (32 classical pipelines, 10 zero-shot LLMs)
(a) Macro-F1
(b) Accuracy
Albertina-900M + SVC-RBF
BERTimbau-large + SVC-RBF
BERTimbau-base + SVC-RBF
Albertina-100M + SVC-RBF
Albertina-900M + LinearSVC
Albertina-900M + LogReg
Albertina-900M + kNN
BERTimbau-large + LinearSVC
BERTimbau-large + LogReg
BERTimbau-base + kNN
Albertina-900M + GradBoost
BERTimbau-large + kNN
BERTimbau-large + GradBoost
Albertina-100M + kNN
Albertina-900M + RandomForest
BERTimbau-large + RandomForest
BERTimbau-base + LinearSVC
BERTimbau-base + LogReg
BERTimbau-base + GradBoost
BERTimbau-base + RandomForest
Albertina-100M + LogReg
Albertina-100M + LinearSVC
Albertina-100M + GradBoost
Albertina-100M + RandomForest
Albertina-900M + DecisionTree
BERTimbau-base + DecisionTree
BERTimbau-large + DecisionTree
Albertina-100M + DecisionTree
Gemma-2 9B
Qwen-2.5 7B
BERTimbau-base + GaussianNB
BERTimbau-large + GaussianNB
Albertina-100M + GaussianNB
Albertina-900M + GaussianNB
Qwen-2 7B
Gemma-3 12B
DeepSeek-R1 8B
Llama-3.1 8B
Phi-3 14B
Qwen-2.5 14B
Llama-3.2 3B
GPT-OSS 20B
0.4
0.5
0.6
0.7
0.8
0.9
1.0
Test macro-F1 (95 % bootstrap CI, 2 000 resamples)
Test accuracy (95 % bootstrap CI)
BERTimbau-base
BERTimbau-large
Albertina-100M
Albertina-900M
Zero-shot LLM