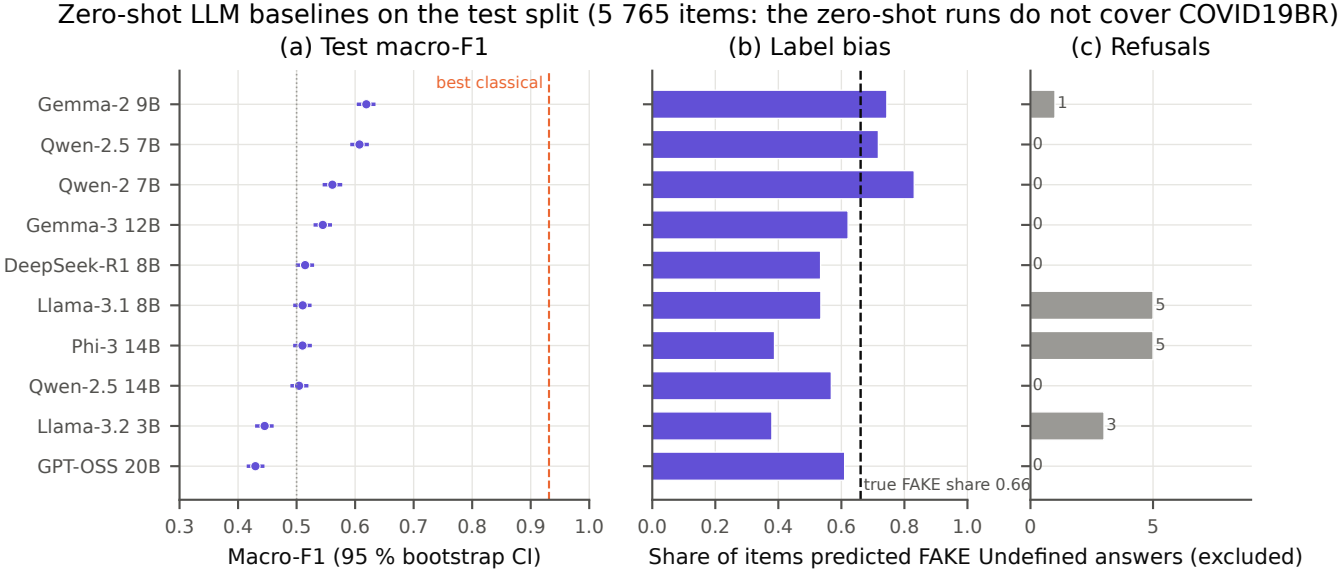
Zero-shot LLM baselines on the test split (5 765 items: the zero-shot runs do not cover COVID19BR)
(a) Test macro-F1
(b) Label bias
(c) Refusals
best classical
Gemma-2 9B
Qwen-2.5 7B
Qwen-2 7B
Gemma-3 12B
DeepSeek-R1 8B
Llama-3.1 8B
Phi-3 14B
Qwen-2.5 14B
Llama-3.2 3B
GPT-OSS 20B
0.3 0.4 0.5 0.6 0.7 0.8 0.9 1.0
Macro-F1 (95 % bootstrap CI)
true FAKE share 0.66
0.0 0.2 0.4 0.6 0.8 1.0
Share of items predicted FAKE
1
0
0
0
0
5
5
0
3
0
0 5
Undefined answers (excluded)